

# Lab 6: RV64 缺页异常处理以及 fork 机制

课程名称：操作系统 指导老师：寿黎但

姓名：庄毅非

## 1、实验目的

- 通过 `vm_area_struct` 数据结构实现对 task 多区域虚拟内存的管理。
- 在 Lab5 实现用户态程序的基础上，添加缺页异常处理 `Page Fault Handler`。
- 为 task 加入 `fork` 机制，能够支持通过 `fork` 创建新的用户态 task。

## 2、实验环境

- Environment in previous lab5

## 3、实验步骤

### 3.1 实现vma

这里主要是对vma机制的实现，首先修改proc.h文件，添加如下的内容。

```
1 // proc.h
2 // 添加 vm_area_struct 结构体
3 struct vm_area_struct {
4     uint64_t vm_start;           /* VMA 对应的用户态虚拟地址的开始 */
5     uint64_t vm_end;             /* VMA 对应的用户态虚拟地址的结束 */
6     uint64_t vm_flags;           /* VMA 对应的 flags */
7     uint64_t vm_content_offset_in_file; /* 对应内容在文件中的偏移量 */
8     uint64_t vm_content_size_in_file; /* 对应内容在文件中的大小 */
9 };
10 // 在 task_struct 中添加 vmas 柔性数组
11 struct task_struct {
12     uint64_t state;
13     uint64_t counter;
14     uint64_t priority;
15     uint64_t pid;
16
17     struct thread_struct thread;
18     pagetable_t pgd;
19
20     uint64_t vma_cnt; // vma的个数
21     struct vm_area_struct vmas[0];
22 };
23 // 添加函数定义
```

```

24 void do_mmap(struct task_struct *task, unsigned long addr, unsigned long length,
    unsigned long flags, unsigned long vm_content_offset_in_file, unsigned long
    vm_content_size_in_file);
25 struct vm_area_struct *find_vma(struct task_struct *task, unsigned long addr);
26 unsigned long clone(struct pt_regs *regs);

```

实现相关功能函数 find\_vma 以及 do\_mmap。

```

1 //proc.c
2 // 寻找addr对应的vma
3 struct vm_area_struct *find_vma(struct task_struct *task, unsigned long addr)
4 {
5     int vma_count = task->vma_cnt;
6     for (int i = 0; i < vma_count; i++)
7     {
8         struct vm_area_struct *move = &(task->vmas[i]);
9         if (addr ≥ move->vm_start && addr < move->vm_end)
10         {
11             return move;
12         }
13     }
14     return NULL;
15 }
16 // 在进程中添加vma
17 void do_mmap(struct task_struct *current_task, unsigned long addr, unsigned long
    length, unsigned long flags, unsigned long vm_content_offset_in_file, unsigned
    long vm_content_size_in_file)
18 {
19     struct vm_area_struct *newvma = &(current_task->vmas[current_task->vma_cnt]);
20     newvma->vm_start = addr;
21     newvma->vm_end = addr + length;
22     newvma->vm_flags = flags;
23     newvma->vm_content_offset_in_file = vm_content_offset_in_file;
24     newvma->vm_content_size_in_file = vm_content_size_in_file;
25     current_task->vma_cnt++;
26 }

```

## 3.2 page fault handle

### 3.2.1 demand paging

修改初始化进程代码，现在在一开始只初始化一个用户进程，并且不使用create\_mapping进行映射，而是使用do\_mmap进行映射。要注意的一点是我在vma中存储的flag是其对应页表中的flag，而不是其右移一位的结果，也就是说如果一个vma对应的页表项的permission是可读可写可执行有效的话，那么其vma对应的flag会是 $(1 \ll 1) \mid (1 \ll 2) \mid (1 \ll 3) \mid 1$ ，而不是 $1 \mid (1 \ll 1) \mid (1 \ll 2)$ ，这样做是为了方便do\_page\_fault函数的实现。对于匿名页的判断，因为只有匿名页对应的vm\_content\_offset\_in\_file和vm\_content\_size\_in\_file能够同时为0，所以我在初始化匿名页的时候会将这两个值设置为0，在处理page\_fault的时候，如果这两个值都是0，那么就是匿名页的映射。

```
1 void task_init()
2 {
3     idle = (struct task_struct *)kalloc();
4     idle->state = TASK_RUNNING;
5     idle->counter = 0;
6     idle->priority = 0;
7     idle->pid = 0;
8     current = idle;
9     task[0] = idle;
10    // 只初始化一个用户进程
11    struct task_struct *launched = task[1] = (struct task_struct *)kalloc();
12    launched->state = TASK_RUNNING;
13    launched->counter = 0;
14    launched->priority = rand();
15    launched->pid = 1;
16    launched->thread.ra = (unsigned long)__dummy;
17    launched->thread.sp = (unsigned long)launched + PGSIZE;
18    launched->thread_info = (struct thread_info *)(alloc_page());
19    launched->thread_info->kernel_sp = launched->thread.sp;
20    launched->thread_info->user_sp = (alloc_page());
21    // 用户页表初始化
22    pagetable_t pgtbl = (unsigned long *)(alloc_page());
23    copy(swapper_pg_dir, pgtbl);
24    // 使用do_mmap分配用户栈这个匿名页
25    do_mmap(launched, USER_END - PGSIZE, PGSIZE, READABLE | WRITABLE | USER |
VALID, 0, 0);
26    // set sstatus
27    unsigned long sstatus = csr_read(sstatus);
28    // set spp bit to 0
29    sstatus &= ~(1 << 8);
30    // set SPIE bit to 1
31    sstatus |= (1 << 5);
32    // set SUM bit to 1
33    sstatus |= (1 << 18);
34    launched->thread.sstatus = sstatus;
```

```

35 // user mode stack
36 launched→thread.sscratch = USER_END;
37
38 // printk("SET [PID = %d] COUNTER = %d\n", launched → pid, launched →
counter);
39 unsigned long satp = csr_read(satp);
40 launched→pgd = (pagetable_t)((((satp >> 44) << 44) | (((unsigned long)(pgtbl)-
PA2VA_OFFSET) >> 12)));
41 load_program(launched);
42 /* YOUR CODE HERE */
43 printk("...proc_init done!\n");
44 }

```

修改load\_program如下,主要也是将create\_mapping修改为do\_mmap.

```

1 static void load_program(struct task_struct *task)
2 {
3     Elf64_Ehdr *ehdr = (Elf64_Ehdr *)uapp_start;
4
5     unsigned long phdr_start = (unsigned long)ehdr + ehdr→e_phoff;
6     int phdr_cnt = ehdr→e_phnum;
7
8     Elf64_Phdr *phdr;
9     for (int i = 0; i < phdr_cnt; i++)
10    {
11        phdr = (Elf64_Phdr *)(phdr_start + sizeof(Elf64_Phdr) * i);
12        if (phdr→p_type == PT_LOAD)
13        {
14            unsigned long va = phdr→p_vaddr;
15            do_mmap(task, PGROUNDDOWN(va), phdr→p_memsz, (phdr→p_flags << 1) |
USER | VALID, phdr→p_offset, phdr→p_filesz);
16        }
17    }
18    task→thread.sepc = ehdr→e_entry;
19 }

```

### 3.2.2 do\_page\_fault

这里是对page\_fault的处理。

首先修改trap.c, 添加对page\_fault的处理, 如果是异常, 并且满足 `scause == PAGE_INSTRUCTION_FAULT` || `scause == PAGE_LOAD_FAULT` || `scause == PAGE_STORE_FAULT`, 那么说明是page\_fault。

```

1 #include "clock.h"
2 #include "printk.h"

```

```

3  #include "proc.h"
4  #include "syscall.h"
5  #include "pagefault.h"
6  void trap_handler(unsigned long scause, unsigned long sepc, struct pt_regs *regs)
7  {
8      if ((scause >> 63) & 1)
9      {
10         if ((scause & 5) == 5)
11         {
12             clock_set_next_event();
13             do_timer();
14         }
15     }
16     else
17     {
18         if (scause == 8)
19         {
20             syscall(regs);
21         }
22         else if (scause == PAGE_INSTRUCTION_FAULT || scause == PAGE_LOAD_FAULT ||
scause == PAGE_STORE_FAULT)
23         {
24             do_page_fault(regs);
25         }
26     }
27 }
28

```

然后是do\_page\_fault函数的实现，这里首先查询stval和cause，然后遍历进程的vma，找到stval所在的vma，如果没找到，那么输出错误信息并返回，如果找到，那么检查其是否满足权限，如果满足权限，那么检查是否是匿名页，如果是匿名页，因为在本实验中只有用户栈是匿名页映射，所以这里就直接将用户栈页进行映射；否则就对对应的磁盘文件进行映射。

```

1  #include "pagefault.h"
2  #include "defs.h"
3  #include "proc.h"
4  #include "vm.h"
5  #include "printk.h"
6  #include "mm.h"
7  extern struct task_struct *current;
8  extern char uapp_start[];
9  void do_page_fault(struct pt_regs *regs)
10 {
11     unsigned long stval = csr_read(stval);
12     unsigned long scause = csr_read(scause);

```

```

13     for (int i = 0; i < current->vma_cnt; i++)
14     {
15         struct vm_area_struct *move = (current->vmas + i);
16         if (move->vm_start ≤ stval && move->vm_end > stval)
17         {
18             // 非法访问
19             if (scause == PAGE_INSTRUCTION_FAULT && ((move->vm_flags & EXECUTABLE)
= 0))
20             {
21                 printk("invalid mem access for PAGE_INSTRUCTION_FAULT\n");
22                 return;
23             }
24             if (scause == PAGE_LOAD_FAULT && ((move->vm_flags & READABLE) == 0))
25             {
26                 printk("invalid mem access for PAGE_LOAD_FAULT\n");
27                 return;
28             }
29             if (scause == PAGE_STORE_FAULT && ((move->vm_flags & WRITABLE) == 0))
30             {
31                 printk("invalid mem access for PAGE_INSTRUCTION_FAULT\n");
32                 return;
33             }
34             pagetable_t pgtbl = (unsigned long *)((((unsigned long)(current-
>pgd)) & 0xffffffff) << 12) + PA2VA_OFFSET);
35             // 检查是否是匿名页
36             if (move->vm_content_offset_in_file == 0 && move-
>vm_content_size_in_file == 0)
37             {
38                 create_mapping(pgtbl, PGROUNDDOWN(move->vm_start), current-
>thread_info->user_sp - PA2VA_OFFSET, (unsigned long)(move->vm_end - move-
>vm_start), move->vm_flags);
39                 return;
40             }
41             else
42             {
43                 // 磁盘文件
44                 unsigned long memPages = (PGROUNDUP(move->vm_end - move-
>vm_start)) / PGSIZE;
45                 // 复制代码段
46                 unsigned long mappedPage = alloc_pages(memPages);
47                 for (int j = 0; j < PGROUNDUP(move->vm_content_size_in_file) / 8;
j++)
48                 {

```

```

49         ((unsigned long *)(mappedPage))[j] = ((unsigned long *)
(PGROUNDDOWN(((unsigned long)(uapp_start + move→vm_content_offset_in_file))))
[j];
50     }
51     // 映射多余内存段
52     create_mapping(pgtbl, move→vm_start, (unsigned long)(mappedPage -
PA2VA_OFFSET), (memPages)*PGSIZE, move→vm_flags);
53     return;
54 }
55 }
56 }
57 printk("[page fault] invalid access\n");
58 return;
59 }

```

### 3.3 fork实现

在syscall.h中添加SYS\_CLONE。

```

1  ...
2      #define SYS_CLONE 220 // fork
3  ...

```

在syscall函数中添加对sys\_clone的处理。

```

1  ...
2  unsigned long sys_clone(struct pt_regs *regs)
3  {
4      return clone(regs);
5  }
6  void syscall(struct pt_regs *regs)
7  {
8      unsigned long sys_code = regs→a7;
9      int return_value = -1;
10     switch (sys_code)
11     {
12         ...
13         case SYS_CLONE:
14         {
15             return_value = sys_clone(regs);
16             break;
17         }
18     }
19     // write result to a0

```

```

20     regs→a0 = return_value;
21     regs→sepc += 4;
22 }

```

clone函数在proc.c中，其调用proc.c中的do\_fork函数。

```

1  unsigned long clone(struct pt_regs *regs)
2  {
3      return do_fork(regs);
4  }
5  unsigned long do_fork(struct pt_regs *regs)
6  {
7      // 进程数++
8      task_count++;
9      int i = task_count;
10     struct task_struct *child = task[i] = (struct task_struct *)kalloc();
11     // 拷贝父进程整个page
12     for (int i = 0; i < PGSIZE / 8; i++)
13     {
14         ((unsigned long *)child)[i] = ((unsigned long *)current)[i];
15     }
16     // 设置子进程相关属性
17     child→counter = 0;
18     child→priority = rand();
19     child→pid = i;
20     // 设置ra为__ret_from_fork函数地址
21     child→thread.ra = (unsigned long)(__ret_from_fork);
22     // 这里设置sp为一个看上去比较奇怪的量，解释如下，事实上父进程在进入trap_handler之前，将
    自己的中断上下文，包括32个寄存器和几个状态寄存器都存储在了地址为current + ((unsigned long)
    (regs)-PGROUNDDOWN((unsigned long)(current)))的位置，由于子进程这时候已经对父进程进行了
    整个page的拷贝，所以子进程可以通过地址(unsigned long)child + (unsigned long)(((unsigned
    long)(regs)-PGROUNDDOWN((unsigned long)(current))))来获取父进程的中断上下文，这里将sp设
    置为它的地址，方便__ret_from_fork的逻辑实现
23     child→thread.sp = (unsigned long)child + (unsigned long)(((unsigned long)
    (regs)-PGROUNDDOWN((unsigned long)(current))));
24     child→thread_info = (struct thread_info *) (alloc_page());
25     child→thread_info→kernel_sp = child→thread.sp;
26     // 拷贝父进程栈
27     child→thread_info→user_sp = (alloc_page());
28     unsigned long parent_user_stack = (unsigned long)(csr_read(sscratch));
29     for (int j = 0; j < PGSIZE / 8; j++)
30     {
31         ((unsigned long *) (child→thread_info→user_sp))[j] = ((unsigned long *)
    (PGROUNDDOWN(parent_user_stack)))[j];
32     }

```



```

33
34     // 用户页表初始化
35     pagetable_t pgtbl = (unsigned long *) (alloc_page());
36     copy(swapper_pg_dir, pgtbl);
37     // 拷贝父进程的mmap
38     child->vma_cnt = current->vma_cnt;
39     for (int j = 0; j < current->vma_cnt; j++)
40     {
41         child->vmass[j] = current->vmass[j];
42     }
43     // 如果有mmap已经映射，那么进行深拷贝
44     for (int i = 0; i < current->vma_cnt; i++)
45     {
46         unsigned long pte = walk(current, current->vmass[i].vm_start);
47         if (pte & VALID)
48         {
49             // 已经mapping，进行拷贝
50             unsigned long *child_pages = (unsigned long *)
(alloc_pages(PGROUNDUP((unsigned long)(current->vmass[i].vm_end) - (unsigned long)
(current->vmass[i].vm_start)) / PGSIZE));
51             for (int j = 0; j < (PGROUNDUP((unsigned long)(current-
>vmass[i].vm_end) - (unsigned long)(current->vmass[i].vm_start)) / 8); j++)
52             {
53                 child_pages[j] = ((unsigned long *) ((unsigned long) current-
>vmass[i].vm_start))[j];
54             }
55             // 执行mapping
56             create_mapping(pgtbl, current->vmass[i].vm_start, (unsigned
long) child_pages - PA2VA_OFFSET, PGROUNDUP((unsigned long)(current-
>vmass[i].vm_end) - (unsigned long)(current->vmass[i].vm_start)), current-
>vmass[i].vm_flags);
57         }
58     }
59     // 修改frame中对应的值
60     struct pt_regs *frame = (struct pt_regs *) ((unsigned long) child + (unsigned
long)(regs) - (unsigned long)(current));
61     frame->a0 = 0;
62     // 修改frame中sp的值为sscratch，事实上就是父进程用户栈的位置
63     frame->sp = csr_read(sscratch);
64     frame->sepc += 4;
65
66     unsigned long sstatus = csr_read(sstatus);
67     child->thread.sstatus = sstatus;
68
69     child->thread.sscratch = child->thread_info->kernel_sp;

```

```

70
71     unsigned long satp = csr_read(satp);
72     child->pgd = (pagetable_t)((((satp >> 44) << 44) | (((unsigned long)(pgtbl)-
    PA2VA_OFFSET) >> 12)));
73     return i;
74 }

```

实现\_\_ret\_from\_fork函数，在entry.S文件中实现。

```

1  .global __ret_from_fork
2  __ret_from_fork:
3      # 这里的sp就是上面说的中断上下文的首地址，子进程从这里对自己的寄存器进行设置。
4      ld t0, 32 * 8(sp)
5      csrw sepc, t0
6      ld t0, 33 * 8(sp)
7      csrw sstatus, t0
8
9      ld x1, 1 * 8(sp)
10     ld x3, 3 * 8(sp)
11     ld x4, 4 * 8(sp)
12     ld x5, 5 * 8(sp)
13     ld x6, 6 * 8(sp)
14     ld x7, 7 * 8(sp)
15     ld x8, 8 * 8(sp)
16     ld x9, 9 * 8(sp)
17     ld x10, 10 * 8(sp)
18     ld x11, 11 * 8(sp)
19     ld x12, 12 * 8(sp)
20     ld x13, 13 * 8(sp)
21     ld x14, 14 * 8(sp)
22     ld x15, 15 * 8(sp)
23     ld x16, 16 * 8(sp)
24     ld x17, 17 * 8(sp)
25     ld x18, 18 * 8(sp)
26     ld x19, 19 * 8(sp)
27     ld x20, 20 * 8(sp)
28     ld x21, 21 * 8(sp)
29     ld x22, 22 * 8(sp)
30     ld x23, 23 * 8(sp)
31     ld x24, 24 * 8(sp)
32     ld x25, 25 * 8(sp)
33     ld x26, 26 * 8(sp)
34     ld x27, 27 * 8(sp)
35     ld x28, 28 * 8(sp)
36     ld x29, 29 * 8(sp)

```

```
37      ld x30, 30 * 8(sp)
38      ld x31, 31 * 8(sp)
39      ld sp, 2 * 8(sp)
40      #-- -- -- -- -
41      sret
```

程序运行效果如下。

```
SET [PID = 4 COUNTER = 9]
switch to [PID = 3 COUNTER = 4]
[U] pid: 3 is running!, global_variable: 201
[U] pid: 3 is running!, global_variable: 202
[U] pid: 3 is running!, global_variable: 203
[U] pid: 3 is running!, global_variable: 204
[U] pid: 3 is running!, global_variable: 205
switch to [PID = 1 COUNTER = 8]
[U] pid: 1 is running!, global_variable: 229
[U] pid: 1 is running!, global_variable: 230
[U] pid: 1 is running!, global_variable: 231
[U] pid: 1 is running!, global_variable: 232
[U] pid: 1 is running!, global_variable: 233
[U] pid: 1 is running!, global_variable: 234
[U] pid: 1 is running!, global_variable: 235
[U] pid: 1 is running!, global_variable: 236
[U] pid: 1 is running!, global_variable: 237
[U] pid: 1 is running!, global_variable: 238
switch to [PID = 4 COUNTER = 9]
[U] pid: 4 is running!, global_variable: 184
[U] pid: 4 is running!, global_variable: 185
[U] pid: 4 is running!, global_variable: 186
[U] pid: 4 is running!, global_variable: 187
[U] pid: 4 is running!, global_variable: 188
[U] pid: 4 is running!, global_variable: 189
[U] pid: 4 is running!, global_variable: 190
[U] pid: 4 is running!, global_variable: 191
[U] pid: 4 is running!, global_variable: 192
[U] pid: 4 is running!, global_variable: 193
[U] pid: 4 is running!, global_variable: 194
[U] pid: 4 is running!, global_variable: 195
switch to [PID = 2 COUNTER = 10]
[U] pid: 2 is running!, global_variable: 243
[U] pid: 2 is running!, global_variable: 244
[U] pid: 2 is running!, global_variable: 245
[U] pid: 2 is running!, global_variable: 246
[U] pid: 2 is running!, global_variable: 247
[U] pid: 2 is running!, global_variable: 248
[U] pid: 2 is running!, global_variable: 249
[U] pid: 2 is running!, global_variable: 250
[U] pid: 2 is running!, global_variable: 251
```

## 4. 思考题

#### 4.1 `uint64_t vm_content_size_in_file;` 对应的文件内容的长度。为什么还需要这个域？

答案：这个问题事实上在lab5的报告中已经回答了，在装载用户程序的时候，lab6使用的是将elf文件加载到内存中，因为在elf文件中的load的段的p\_memsz和p\_filesz一般并不相等，一个load段可能有无需初始化的变量或者数组等数据，这些数据都会在程序装载的时候初始化为0，如果在elf文件中为这些数据预留空间的话，就会造成磁盘空间的浪费。但是这些数据在内存中是需要对应的空间来保障程序运行的，所以在许多情况下p\_memsz会稍大于p\_filesz。所以我们必须分别记录每个段的文件中的大小和内存中的起止位置，才能够进行page\_fault的处理。

#### 4.2 `struct vm_area_struct vmas[0];` 为什么可以开大小为0的数组？这个定义可以和前面的vma\_cnt换个位置吗？

答案：这是gcc支持的一个特性，叫做柔性数组，允许我们在运行的时候通过申请不同的内存大小来动态开辟数组，这个定义不能和vma\_cnt交换，这是柔性数组本身规定的，其必须在结构的最后。

#### 4.3 想想为什么只要拷贝那些已经分配并映射的页，那些本来应该被分配并映射，但是暂时还没有因为Page Fault而被分配并映射的页怎么办？

答案：对于那些父进程已经分配和映射的页面，很有可能父进程已经对其进行了修改等写入行为，如果我们在fork的时候不进行拷贝的话，那么子进程在返回用户区间运行的时候，就可能因为读取到和父进程不同的值，导致运行出错，一个例子就是父进程在0x10200写入了一个32（原来是-1），然后在fork之后通过读取这个32来执行程序，子进程如果不进行拷贝，从磁盘上拷贝0x10200对应的vma，之后读取到了这个-1，那么本来子进程应该和父进程读取到一样的值的，这里却不一样了，这就可能导致程序运行出错。

对于那些还没被映射的页，由于父进程没有通过do\_page\_fault进行映射，那么其肯定也没有修改对应page中的值，子进程也不需要拷贝，直接从磁盘中读取即可。

#### 4.4 参考task\_init 创建一个新的task，将的parent task的整个页复制到新创建的task\_struct页上，这一步复制了哪些东西？

答案：事实上这一步是对整个task\_struct的复制，从task\_struct的结构定义可以看出，这里主要是复制了线程的ra、sp、12个s寄存器，sepc寄存器，sstatus寄存器，sscratch寄存器，以及父进程的vma。当然也复制了父进程状态，计数器，页表等。

#### 4.5 将thread.ra 设置为 `__ret_from_fork`，并正确设置 `thread.sp`。仔细想想，这个应该设置成什么值？可以根据child task的返回路径来倒推。

答案：这里我将sp设置为 `(unsigned long)child + (unsigned long)((((unsigned long)(regs) - PGROUNDDOWN((unsigned long)(current)))))`，事实上，regs是父进程存储中断上下文的地址，由于child在此前已经对父进程的整个task\_struct结构体进行了拷贝，所以这里 `(unsigned long)child + (unsigned long)((((unsigned long)(regs) - PGROUNDDOWN((unsigned long)(current)))))` 指向的就是子进程中存储父进程中中断上下文的地方，我们后续可以通过这个地址，在\_\_ret\_from\_fork中进行子进程状态的恢复。

4.6 利用参数 `regs` 来计算出 `child task` 的对应的 `pt_regs` 的地址, 并将其中的 `a0`, `sp`, `sepc` 设置成正确的值。为什么还要设置 `sp`?

答案: 这里的`sp`事实上需要是父进程用户栈的虚拟地址, 如果不设置的话, 那么存储的会是父进程`kernel`栈在存储中断上下文时的地址, 导致程序运行出错。在我的实现中, 是在`fork`的逻辑里将`sp`设置为`csr_read(sscratch)`, 也就是父进程栈的地址。

## 5. 更多测试用例

使用了实验手册中提供的斐波那契函数进行测试, 程序运行无误。

```
SET [PID = 1 COUNTER = 10]
SET [PID = 2 COUNTER = 5]
[U-CHILD] pid: 2 is running! the 38th fibonacci number is 39088169 and the number @ 962 in the large array is 962
switch to [PID = 1 COUNTER = 10]
[U-PARENT] pid: 1 is running! the 37th fibonacci number is 24157817 and the number @ 963 in the large array is 963
[U-PARENT] pid: 1 is running! the 38th fibonacci number is 39088169 and the number @ 962 in the large array is 962
SET [PID = 1 COUNTER = 2]
SET [PID = 2 COUNTER = 9]
[U-PARENT] pid: 1 is running! the 39th fibonacci number is 63245986 and the number @ 961 in the large array is 961
switch to [PID = 2 COUNTER = 9]
[U-CHILD] pid: 2 is running! the 39th fibonacci number is 63245986 and the number @ 961 in the large array is 961
SET [PID = 1 COUNTER = 4]
SET [PID = 2 COUNTER = 4]
switch to [PID = 1 COUNTER = 4]
switch to [PID = 2 COUNTER = 4]
[U-CHILD] pid: 2 is running! the 40th fibonacci number is 102334155 and the number @ 960 in the large array is 960
SET [PID = 1 COUNTER = 10]
SET [PID = 2 COUNTER = 5]
switch to [PID = 1 COUNTER = 10]
[U-PARENT] pid: 1 is running! the 40th fibonacci number is 102334155 and the number @ 960 in the large array is 960
SET [PID = 1 COUNTER = 10]
SET [PID = 2 COUNTER = 4]
switch to [PID = 2 COUNTER = 4]
switch to [PID = 1 COUNTER = 10]
[U-PARENT] pid: 1 is running! the 41th fibonacci number is 165580141 and the number @ 959 in the large array is 959
SET [PID = 1 COUNTER = 7]
SET [PID = 2 COUNTER = 5]
switch to [PID = 2 COUNTER = 5]
switch to [PID = 1 COUNTER = 7]
SET [PID = 1 COUNTER = 8]
SET [PID = 2 COUNTER = 8]
switch to [PID = 2 COUNTER = 8]
[U-CHILD] pid: 2 is running! the 41th fibonacci number is 165580141 and the number @ 959 in the large array is 959
```