

db hw10

3200105872 庄毅非

15.2 Consider the bank database of Figure 15.14, where the primary keys are underlined, and the following SQL query:

```
1 select T.branch\_name
2 from branch T, branch S
3 where T.assets > S.assets and S.branch_city = "Brooklyn"
```

Write an efficient relational-algebra expression that is equivalent to this query. Justify your choice.

Answer:

$$\Pi_{T.branch_name}(\Pi_{(assets,branch_name)}(\rho_T(branch)) \bowtie_{(T.assets > S.assets)} (\Pi_{assets}(\sigma_{branch_city='Brooklyn'}(\rho_S(branch)))))$$

In this way, the two operation sets for performing the join operation are minimized, and the amount of data passed during the operation is reduced, and the simpler selection operation is performed first, reducing the cost of subsequent projection and join operations. In the end, the total cost of all operations is minimized.

15.3 Let relations $r_1(A, B, C)$ and $r_2(C, D, E)$ have the following properties: r_1 has 20,000 tuples, r_2 has 45,000 tuples, 25 tuples of r_1 fit on one block, and 30 tuples of r_2 fit on one block. Estimate the number of block transfers and seeks required using each of the following join strategies for $r_1 \bowtie r_2$:

1. Nested-loop join.

Answer: Firstly, the number of disk blocks occupied by r_1 is $b_1 = 20000 / 25 = 800$, and the number of disk blocks occupied by r_2 is $b_2 = 45000 / 30 = 1500$.

If buffer size $M > 800$, then we need $800 + 1500 = 2300$ disk accesses. Else the size of buffer is less or equal to 800. Assume r_1 is the outer relation, then our accesses to the disk is $800 + 20000 * 1500 = 30000800$. If r_2 is the outer relation, then our accesses to the disk is $1500 + 800 * 45000 = 36001500$.

b. Block nested-loop join:

Answer:

If buffer size $M > 800$, then we need $800 + 1500 = 2300$ disk accesses.

Else if r_1 is the outer relation, we need $\lceil 800/(M-1) \rceil * 1500 + 800$ disk accesses.

Else if r_2 is the outer relation, we need $\lceil 1500/(M-1) \rceil * 800 + 1500$ disk accesses.

c. Merge-join:

Answer:

If buffer size $M > 800$, then we need $800 + 1500 = 2300$ disk accesses.

If r_1 and r_2 are sorted, we need $800 + 1500$ disk accesses.

Else we need firstly sort r_1 and r_2 , and we need write back to disk, whose cost is $2 * 800(\lceil \log_{M-1}(800/M) \rceil + 1) + 2 * 1500(\lceil \log_{M-1}(1500/M) \rceil + 1)$ disk accesses. Then we need another $1500 + 800 = 2300$ disk accesses. So our final result is $2 * 800(\lceil \log_{M-1}(800/M) \rceil + 1) + 2 * 1500(\lceil \log_{M-1}(1500/M) \rceil + 1) + 2300$

4. Hash join.

Answer: Assume r_1 is the build relation and r_2 is the probe relation. If $M > 800 / M + 1$, then we don't need to recursively partition. The number of access of disk is $3 * (800 + 1500) = 6900$, else number of access of disk is $2 * (800 + 1500) * (\lceil \log_{M-1}(800) \rceil - 1) + 800 + 1500 = 4600 * (\lceil \log_{M-1}(800) \rceil - 1) + 2300$